

Technical Report

1. Introduction

Healthcare fraud costs the U.S. healthcare system more than \$68 billion annually, draining resources from legitimate patients. Since CMS (Centers for Medicare & Medicaid Services) can manually investigate only a small fraction of suspicious cases, there is a strong need for an intelligent automated detection system.

This project develops an **end-to-end machine learning pipeline** capable of identifying high-risk healthcare providers using a real Medicare fraud dataset from Kaggle.

The goals are:

- Detect fraudulent providers at the provider level
- Handle severe class imbalance (~9% fraud)
- Build interpretable and practical models
- Justify all steps: data cleaning, feature engineering, model selection, tuning, and error analysis

All implementation steps were originally completed in a **single notebook (mlproj2)** before being divided into the required three-notebook structure. Because of this workflow, some steps may depend on earlier transformations.

2. Dataset Description

The dataset contains four CSV files at different granularities:

1. Train_Beneficiarydata.csv

- Beneficiary demographics (DOB, DOD, Gender, Race, State)
- Chronic condition indicators (e.g., CHF, cancer, diabetes)
- **Granularity:** BeneID

2. Train_Inpatientdata.csv

- Inpatient hospital claims
- Claim dates, reimbursement amounts, deductibles
- Diagnosis codes, physician IDs
- **Granularity:** Claim-level (BeneID → Provider)

3. Train_Outpatientdata.csv

- Outpatient visits, tests, procedures
- Similar structure to inpatient
- **Granularity:** Claim-level

4. Train_Labels.csv

- Provider fraud labels ("Yes" / "No")
- **Granularity:** Provider

Key Relationships

- **BeneID** links beneficiary → claims
- **Provider** links claims → fraud label

Thus, modeling must be done at the provider level, requiring extensive aggregation across all tables.

3. Data Understanding & Exploration (1.5.1)

3.1 Initial Inspection

We performed:

- `.info()` for data types
- `.isnull().sum()` for missing values
- `.shape + .nunique()`
- Validated date columns
- Checked Beneficiary–Claim–Provider coverage

Findings:

- Some missing dates
- Missing chronic conditions
- Strongly skewed reimbursement amounts
- Providers appear across datasets inconsistently

3.2 Beneficiary Analysis

We converted DOB/DOD to datetime and computed:

- Age distribution (mainly 70–80+)
- Gender distribution
- Race distribution
- Renal disease prevalence
- Chronic condition prevalence
- State distribution

Beneficiaries exhibit many chronic conditions, typical for Medicare.

3.3 Claims Analysis

Performed on both inpatient & outpatient claims:

- Monthly claim counts

- Claim duration
- Reimbursement and deductible distributions
- Temporal trends
- Outliers
- Geographic patterns

Findings:

- Monthly claim volume fluctuates
 - Reimbursement amounts highly skewed
 - State-level concentration of claims
-

4. Provider-Level Aggregation Strategy

Since labels are provider-level, we aggregated all claim-level features.

4.1 Inpatient Aggregations

For each provider:

- Sum / mean / std of **InscClaimAmtReimbursed**
- Sum / mean of **DeductibleAmtPaid**
- Count of inpatient claims
- Unique attending / operating / other physicians

4.2 Outpatient Aggregations

Identical aggregator set applied.

4.3 Combined Provider Features

Created:

- `total_claims`
- `inpatient_ratio`
- `avg_claim_amount`
- `physician_variety`
- Chronic condition percentages per provider

This produced the provider-level feature table for modeling with **51 engineered features**.

5. Advanced Feature Engineering

5.1 High-Cost Claim Percentages

Using 90th-percentile thresholds:

- pct_high_cost_inpatient
- pct_high_cost_outpatient
- pct_high_cost_total

5.2 Chronic Condition Intensity

Converted chronic condition indicators to binary and computed:

- Mean prevalence per provider
- Chronic condition ratios for inpatient/outpatient
- pct_chronic_patients

5.3 Operational Features

- 30-day readmission rate
 - Physician utilization
 - Cost-per-physician
-

6. Class Imbalance Analysis (1.5.2)

Fraud distribution:

- **No fraud:** ~91.13% (3,816 providers)
- **Fraud:** ~8.87% (399 providers)

Approximately **9.5:1 imbalance ratio.**

Visualized:

- Fraud vs. non-fraud pie chart
- Class counts

Insight:

A model predicting “No fraud” always would still achieve 91.13% accuracy.

7. Imbalance Strategy: Class Weighting

Class weights using sklearn's `compute_class_weight('balanced')`:

- **Class 0 weight (No fraud):** 0.567
- **Class 1 weight (Fraud):** 4.258

Why class weighting?

1. No data loss
2. Avoids synthetic samples
3. Penalizes false negatives
4. Works with many algorithms

For XGBoost: `scale_pos_weight = 7.52.`

8. Algorithm Selection (1.5.3)

Models implemented with tuning:

1. Decision Tree
2. Random Forest
3. Gradient Boosting / XGBoost
4. Logistic Regression

SVM Considered but Not Implemented

Due to:

- High computational cost
 - Poor performance on imbalanced tabular data
 - Low interpretability
 - Outperformed by tree ensembles + logistic regression
-

9. Validation Strategy

We used a robust two-part strategy:

✓ **5-fold cross-validation** (GridSearchCV)

✓ **20% hold-out test set**

Data splits:

- Total providers: 4,215
 - Training: 3,372
 - Test: 843
-

10. Comparative Model Performance

Baseline Models (class weighting):

Model	Precision	Recall	F1	Accuracy
Decision Tree	0.444	0.449	0.447	0.859
Random Forest	0.667	0.430	0.523	0.900
XGBoost	0.553	0.533	0.543	0.886
Logistic Regression	0.413	0.794	0.543	0.830

Tuned Models:

Model	Precision	Recall	F1	Accuracy
Decision Tree (Tuned)	0.397	0.832	0.538	0.819
Random Forest (Tuned)	0.593	0.654	0.622	0.899
XGBoost (Tuned)	0.464	0.841	0.598	0.856
Logistic Regression (Tuned)	0.449	0.860	0.590	0.848

Key Findings

- Recall most important
- Logistic Regression balances precision/recall best
- Trees yield high precision but miss more fraud
- Tuning dramatically boosts recall

11. Experiment Log — Detailed Trial Documentation

11.1–11.8 Model Summaries

- **Baseline Decision Tree:** Recall 0.449, unstable
- **Tuned Decision Tree:** Recall 0.832
- **Baseline Random Forest:** High precision, low recall
- **Tuned Random Forest:** Recall 0.654
- **Baseline XGBoost:** Moderate
- **Tuned XGBoost:** Recall 0.841
- **Baseline Logistic Regression:** Strong recall 0.794
- **Tuned Logistic Regression:** Best overall recall 0.860

12. Final Model Selection: Logistic Regression

Primary Reason

- **Highest recall (0.860)**

- Minimizes false negatives
- Aligns with CMS mission

Secondary Advantages

- Interpretability
- Stability
- Efficiency
- Probabilistic outputs

Top Fraud Indicators

- High % expensive claims
 - Many unique physicians
 - High chronic condition intensity
 - Balanced claim mix correlated negatively with fraud
-

13. Evaluation & Error Analysis

Test Set Results (Tuned Logistic Regression):

- Accuracy: 0.848
- Precision: 0.449
- Recall: 0.860
- F1: 0.590

Confusion Matrix

		Predicted	
		No	Yes
Actual	No	[TN=625]	[FP=111]
	Yes	[FN=15]	[TP=92]

Error Patterns

False Positives (111):

- High claim amounts
- Multiple physicians
- Explosive costs

False Negatives (15):

- Low claim volume
- Subtle fraud

- Hard to distinguish
-

14. Business Impact Analysis

With Logistic Regression

- Fraud detection: **86%**
- FP rate: **15.1%**
- Investigation workload: **24%**
- Missed fraud: **15 providers**

Compared with Other Models

- Random Forest: misses 37 fraud cases
 - XGBoost: high recall but less interpretable
 - Decision Tree: unstable performance
-

15. Limitations & Future Work

Improvements

- Threshold tuning
- Cost-sensitive learning
- Ensemble combinations

Advanced Features

- Temporal patterns
- Network analysis
- Anomaly detection

Operational

- Model monitoring
 - Feedback loop retraining
 - Explainability dashboard
-

16. Conclusion

This project delivered a **complete Medicare fraud detection pipeline**, including:

1. Comprehensive data exploration

2. 51 engineered provider-level features
3. Class-weighting for imbalance
4. Full model comparison
5. Logistic Regression with **0.860 recall** selected
6. Detailed error analysis

The final system is **transparent, practical, and aligned with CMS objectives**, maximizing fraud detection while supporting investigators.

Deployment Recommendation:

Use Logistic Regression with a **0.3–0.4 probability threshold** and integrate an explainability dashboard.